

NOODLES V0.1: PROVISIONAL SPECIFICATION

NICHOLAS BRUNHART-LUPO

CONTENTS

1. Introduction	2
2. Rationale & Design Goals	2
3. Architecture	2
3.1. Communication	3
4. Concepts	3
4.1. Document	3
4.2. Identifiers	3
4.3. Objects	5
4.4. Tables	5
4.5. Signals and Methods	5
4.6. Buffers	5
4.7. Mesh	5
4.8. Materials	5
4.9. Textures	5
4.10. Lights	6
5. Common Message Elements	6
5.1. Any Type	6
6. Server Messages	7
6.1. Root Message	7
6.2. Objects	8
6.3. Tables	9
6.4. Buffers	10
6.5. Geometry	10
6.6. Texture	11
6.7. Material	11
6.8. Lights	12
6.9. Signals & Methods	12
6.10. Signal Invoke & Method Reply	13
6.11. Document	13
7. Client Messages	14
7.1. Root Message	14
7.2. Introduction	14

7.3. Method Invocation	14
7.4. Asset Refresh	15
8. Semantics	15
8.1. Tables	15
8.2. Objects	18
9. Operation & Lifecycle	20
9.1. Websocket Messages	20
9.2. Connection	20
Appendix A. Common Flatbuffer Specification	21
Appendix B. Server Message Flatbuffer Specification	23
Appendix C. Client Message Flatbuffer Specification	29

1. INTRODUCTION

This document entails a specification for a distributed scene-graph wireline protocol suitable as a substrate for shared interactive visualizations. It also lays out concepts for the supporting implementations that would provide such visualizations.

2. RATIONALE & DESIGN GOALS

- The intent of this document is simplicity, to get a working version implemented so that further improvements can be identified.
- The structure here is not intended to mirror the use-case of the HTML DOM + Javascript where code is shipped to clients. That would be restrictive, as it requires the clients either interpret or compile and run code on command. Some clients, such as integrated head mounted systems, do not allow compilation, or are not sufficient computing platforms.
- Trying to mirror just the HTML DOM part has issues as well; a number of 3D declarative implementations (like QML 3D), all operate on a scenegraph under the hood. It seems more fruitful to just target the scenegraph for modification, and perhaps (as part of the server library) have a declarative component there.
- A shared document is desired here, as opposed to the standard browser case where every client has their own copy of state.
- Code listings are provided as an example and for exposition only. Clients and servers may be written in any language as long as they conform to the proper wireline protocol.

3. ARCHITECTURE

The system envisions the use of four components, two of which fall under this specification. The Server Library presents a visualization to one or more connected clients through a synchronized scenegraph. Client requests and messages are passed on for handling to the

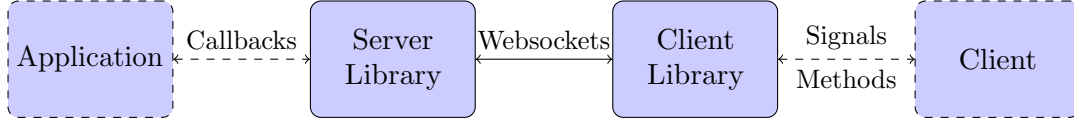


FIGURE 1. System architecture. Note that there may be more than one client. Elements with solid lines fall under this specification.

application code, which can manipulate the scenegraph in response. These changes are then published and sent to clients.

The Client Library connects to a server, and maintains the synchronized scenegraph. This scenegraph is query-able by the client. Clients then can interpret and present the scenegraph to the user in the way they see fit. For example, an immersive graphics engine client can draw the scenegraph as is, while a 2D client can choose to present only a subset of the graph. A command line (i.e. Python) client may ignore the scenegraph completely to merely make use of the messaging and method invocation functionality. This also allows each client to customize the interactions available in a way that best aligns with their form factor.

3.1. Communication. Communication between the libraries is achieved over Websocket connections. All messages are sent over the binary channel of the WebSocket using Flatbuffers.

Client-to-client notification is not supported, and must first pass through the server.

The bulk of communication is from server to client.

This spec is intended to be implemented in a secure network, with the presumption that those that connect to the server are trusted. Provision for security will come later, as is the case with everything, because security is hard and makes my brain bleed.

3.1.1. Flatbuffers. For performance reasons, the *in-situ* capabilities of the serialization medium down-selected available options to Flatbuffers and Cap'n'proto. Both were explored. Table 1 compares the two in rough terms. In the end Flatbuffers won out due to more language support out of the box.

4. CONCEPTS

The objective of the system is to synchronize, as best as possible, the document between the client and the server. This is accomplished through the use of discrete messages.

4.1. Document. The Document represents the visualization. It is an entity-component model, with an Object as the core entity, and Tables being a secondary entity.

The document is implicit. The other elements are explicit.

4.2. Identifiers. Identifiers are a pair of 32-bit unsigned integers; the first being a slot number, and the second being a generation count. This allows non-hashed storage, as there should be no two elements with the same slot number, so it can be used as an index in an

	Pro	Con
Cap'n'Proto	Created by Protobuf developers, strong pedigree. Excellent JSON inter-operation.	More complex internal formats. Fewer languages supported out of the box. Some packages for other languages are of lower quality. Default serialization code has performance issues. ¹
Flatbuffers	More languages supported out of the box. Simple internal format, more format features (such as maps, field deprecation, evolution). Easier to obtain performance increases using existing serialization code.	API for some languages is horrible. Some languages require schema to have specific design, adding indirection.

TABLE 1. Serialization Format Comparison

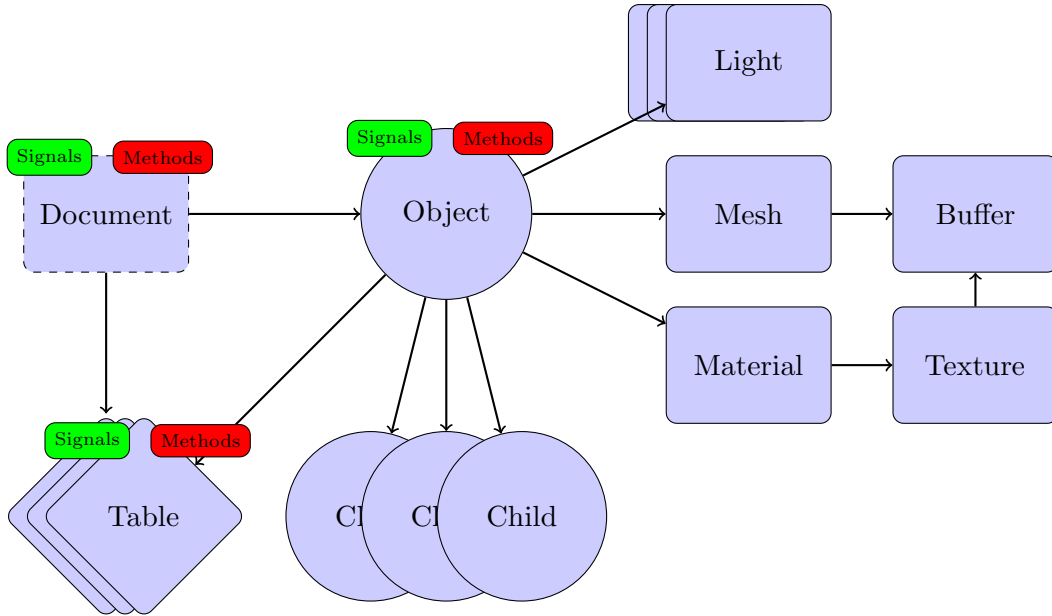


FIGURE 2. Document structure.

array. The generation number is used to help identify if a slot has been recycled by the server, and thus allow detection of stale identifier use.

An identifier where either the slot and generation are the maximum unsigned integer value is the ‘null’ ID.

4.3. Objects. Each object is provided with an Object ID. Objects are rendered in a hierarchy, starting from a root object with the ID 0. Objects can have any number of children.

Each object is a possibly render-able object, and has a transformation, an optional name, a parent Object ID, a mesh (what to draw), a material (how to draw it), a number of lights, and links to tables. Objects also have a set of string tags, and attached methods and signals. Objects can also be instance rendered.

Objects are mutable.

4.4. Tables. Tables are a structured way to transmit row oriented data. They consist of a header (list of column names), and rows. Attached signals and methods are used to allow clients to modify the data in the table or fetch records (but only when first subscribed to).

4.5. Signals and Methods. Signals are notifications from the server to the client. They may contain data, and may come from the document, objects, or tables.

Methods are requests to the server from the client. They may take a set of data parameters, and they may return data as well. They must have a contextual object that they are called on, otherwise they are called on the Document. During the course of a method invocation, signals from the server could be generated.

Each method invocation is tracked by a client-generated arbitrary string. These shall be unique and never re-used. For servers, every method must generate a reply message; the only exception is if the client did not provide an invocation identifier string.

There is a possibility that a method could be called on an object, that is then subsequently deleted, or replaced. In this case, a reply is still generated, and not squashed by the server. Thus a client should be able to handle replies on objects that no longer exist.

Methods and signals are immutable.

4.6. Buffers. A buffer is an opaque block of bytes. This allows for efficient storage and transfer of large assets. These assets can be sent either inline through the WebSocket, or can be supplied through a URL that the client can fetch the buffer from.

Buffers are immutable and referenced from meshes and textures.

4.7. Mesh. Meshes define the geometry that is to be rendered. They consist of references to a buffer for number of components (see Table 2).

Meshes are mutable.

4.8. Materials. This should be a PBR based material, featuring basic elements: base color, metallic, roughness, including an optional texture for base colors. The material only applies to the node it is attached to. Note that though the material is specified in PBR, the client may use Phong or other interpretations of the specified material in order to meet performance goals. The material may also specify that blending should be used; the blending function is src_α and $1 - src_\alpha$.

Materials are mutable.

4.9. Textures. Textures reference images (in Buffers) to be used by a material. Textures are mutable.

Component	Type	Value Type	Count
Position	Vertex	float	3
Normals	Vertex	float	3
TexCoords	Vertex	unsigned short	2
Colors	Vertex	unsigned byte	4
Lines	Index	unsigned short	2
Triangles	Index	unsigned short	3

TABLE 2. Mesh components

4.10. **Lights.** Lights describe illumination sources. They are mutable.

5. COMMON MESSAGE ELEMENTS

This section discusses common elements to both the server and client message portions of the specification.

5.1. **Any Type.** The any type is the foundation value type. it is composed as follows:

LISTING 1. Any Definition

```

1  union AnyIDType {
2      ObjectID,
3      TableID,
4      SignalID,
5      MethodID,
6      MaterialID,
7      GeometryID,
8      LightID,
9      TextureID,
10     BufferID
11 }
12
13 table AnyID {
14     id : AnyIDType (required);
15 }
16
17 table MapEntry {
18     name : string (key);
19     value : Any;
20 }
21
22 table Text { text : string; }
23 table Integer { integer : int64; }
24 table IntegerList { integers : [int64]; }
25 table Real { real : double; }
26 table RealList { reals : [double]; }
27 table Data { data : [byte]; }
28 table AnyList { list : [Any]; }

```

```

29 table AnyMap { entries : [MapEntry]; }
30
31 union AnyType {
32     Text,
33     Integer,
34     IntegerList,
35     Real,
36     RealList,
37     Data,
38     AnyList,
39     AnyMap,
40     AnyID
41 }

```

It is a discriminated union of text, integers, real, or bytes. It also contains a generic list, and string-keyed map. For efficiency, there is also a real-list and integer-list type which allow contiguous storage or access of those elements.

In method and signal API terms, real-lists and integer-lists are coerce-able. That means, for example, that a list of reals may be provided by the `RealList` type, or as a `AnyList` of `Real`.

6. SERVER MESSAGES

Here we discuss the messages that are sent from the server to the client.

Almost all components have strict lifetimes defined by creation and deletion messages. Some messages are also used to update an existing component. Therefore, if a create-update message is received by the client for a component/entity of an ID that it has never seen before, that is the creation milestone. Otherwise, it is an update message. Update messages are treated differently; presence of a key in the update message table implies that the new value should overwrite a previously received value for that key.

6.1. Root Message. The server sends messages by the root type `ServerMessages`. This is an array of a union of creation/deletion messages.

LISTING 2. Server Root Messages

```

1 union ServerMessageType {
2     MethodCreate,
3     MethodDelete,
4     SignalCreate,
5     SignalDelete,
6     ObjectCreateUpdate,
7     ObjectDelete,
8     BufferCreate,
9     BufferDelete,
10    MaterialCreateUpdate,
11    MaterialDelete,
12    TextureCreateUpdate,
13    TextureDelete,
14    LightCreateUpdate,

```

```

15     LightDelete ,
16     GeometryCreate ,
17     GeometryDelete ,
18     TableCreateUpdate ,
19     TableDelete ,
20     DocumentUpdate ,
21     DocumentReset ,
22     SignalInvoke ,
23     MethodReply
24 }
25
26 table ServerMessage {
27     message : ServerMessageType;
28 }
29
30 table ServerMessages {
31     messages : [ ServerMessage ];
32 }
33
34 root_type ServerMessages;

```

6.2. **Objects.** Objects are created, updated and destroyed in Listing 3.

LISTING 3. Object Messages

```

1  table TextDefinition { // Text plane, normal +z, up is +y, center: obj
   origin
2      text      : string (required); // String to render
3      font      : string (required); // Approximate font to use (e.g. Arial)
4      height    : float; // The height of the text. The width will be
5      opt_width : float; // Optional width of font.
6  }
7
8  // non-atomic update
9  table ObjectCreateUpdate {
10     id          : ObjectID (required);
11     name        : string;
12     parent      : ObjectID;
13     transform    : Mat4;
14     material     : MaterialID;
15     mesh        : GeometryID;
16     lights      : [LightID];
17     tables      : [TableID];
18     instances    : [Mat4];
19     tags        : [string];
20     methods_list : [MethodID];
21     signals_list : [SignalID]; // Dont use "signals" to avoid Qt conflict.
22     text        : TextDefinition;
23 }
24
25 table ObjectDelete {

```



```

26 |     id      : ObjectID (required);
27 | }

```

Each element is optional, with the exception of the object id.

If the object ID has not been seen before by the client, it is assumed to be a new object. If there is no transform in the message, it is assumed to be the identity.

If the object ID has been seen before and not deleted, it should update the existing object with the elements that are provided in the message. For example, an message for Object ID 5 that contains a transform will only update object 5's transform, and not change other elements. In another example, to detach a material from an object, an update message with a null material ID is used.

Instances of the underlying meshes are specified by a list of matrices, with a matrix per instance. Using column major ordering, Matrix 1, shows how position p , rotation r (as a quaternion), color c , and scale s are specified. Access is assumed to be by column, i.e. $M_0 = p$.

$$(1) \quad \begin{pmatrix} p_x & c_r & r_x & s_x \\ p_y & c_g & r_y & s_y \\ p_z & c_b & r_z & s_z \\ 0 & 0 & r_w & 0 \end{pmatrix}$$

Text can be added to an object by means of the optional **TextDefinition** part of the message. Text is to be rendered how the client sees fit, with the orientation to be centered at the object, the text perpendicular to $+Z$ and up being $+Y$. The height of the text must be specified; the text will then automatically use the font information to compute the text width. If the optional width is specified, then the text shall, keeping the proper font aspect ratio, try to fill the bounds provided.

6.3. Tables. Tables are created and destroyed in the following messages.

LISTING 4. Table Messages

```

1 | table TableCreateUpdate {
2 |     id      : TableID (required);
3 |     name    : string;
4 |     meta    : string;
5 |     methods_list : [ MethodID ];
6 |     signals_list : [ SignalID ];
7 | }
8 |
9 | table TableDelete {
10 |     id      : TableID (required);
11 | }

```

Tables have names (which shall be unique), a metadata JSON string, methods, and signals.

6.4. **Buffers.** Buffers are created and destroyed in the following messages.

LISTING 5. Buffer Messages

```

1  table BufferCreate {
2      id      : BufferID (required);
3      bytes   : [byte]; // This could be empty
4      url     : string; // This could be empty too
5      url_size : uint64; // If you are passing a URL, this should be set
6
7      // If both are empty, data is coming, just hang on for another message
7      // of
8      // this kind with the bytes. you may have to request a refresh message
9  }
10
11 table BufferDelete {
12     id      : BufferID (required);
13 }

```

Buffers are either inline (in the `bytes` field) or provided as a URL. If a URL is supplied the size of the buffer must be passed as well. If neither is supplied, the server has the data inline, but has deemed it too large to send immediately to avoid stalling clients. In this case, the server would do well to supply the data through another port and use the URL feature, but some servers are unable to do this. In this case, where both the bytes and URL feature are empty, the `url_size` field must still be filled for client pre-allocation. At intervals, the client can send a refresh message to fill in the missing buffers to avoid burdening the websocket (see Section 7.4).

6.5. **Geometry.** Geometries are created, and destroyed in the following messages.

LISTING 6. Geometry Messages

```

1  table ComponentRef {
2      id      : BufferID (required);
3      start   : uint64;
4      size    : uint64;
5      stride  : uint64;
6  }
7
8  table GeometryCreate {
9      id      : GeometryID (required);
10
11     min_extent : Vec3;
12     max_extent : Vec3;
13
14     positions  : ComponentRef; // Vec3
15     normals    : ComponentRef; // Vec3
16     texCoords  : ComponentRef; // Vec2
17     colors     : ComponentRef; // U8Vec4
18     lines      : ComponentRef; // U16Vec2
19     triangles  : ComponentRef; // U16Vec3

```

```

20 }
21
22 table GeometryDelete {
23     id : GeometryID (required);
24 }

```

Limitations in Flatbuffer IDL require some notes here. Geometries are defined by ranges of a buffer for their components. These ranges, in the `ComponentRef` type, require a buffer, and also *require* a start byte offset of the buffer, as well as a byte size field, and the byte stride between vertex elements.

The `min_extent` and `max_extent` fields are required so that clients can efficiently determine culling boundaries.

The position field is required. The other vertex components (`normals`, `texCoords`, `colors`) are optional, but recommended. If there is no normal, the mesh should be rendered without lighting to avoid graphical artifacts. If there are no texture coordinates, the coordinate (0,0) should be assumed for each vertex. If there are no colors, the color (1,1,1,1) should be assumed for each vertex.

Index elements are specified in `lines` and `triangles`. Only one of these should be active. These specify the indices for line segments and triangles respectively. The stride for these components must be zero, i.e. they must be tightly packed.

6.6. Texture. Textures are created, updated, and destroyed in the following messages.

LISTING 7. Texture Messages

```

1 table TextureCreateUpdate {
2     id          : TextureID (required);
3     reference   : BufferRef (required);
4 }
5
6 table TextureDelete {
7     id          : TextureID (required);
8 }

```

Textures, specify a buffer range for an image. For portability, these are to be in the on-disk format for PNG, JPG, or EXR. KTX is also allowed, but the user should be aware that not all clients can support it.

6.7. Material. Materials are created, updated, and destroyed in the following messages.

LISTING 8. Material Messages

```

1 table MaterialCreateUpdate {
2     id          : MaterialID (required);
3     color       : Vec4;
4     metallic    : float;
5     roughness   : float;
6     use_blending : bool;
7     texture_id  : TextureID;
8 }

```

```

9
10 table MaterialDelete {
11     id          : MaterialID (required);
12 }

```

The `color` key defines colors in the range of 0 – 1 for r, g, b, a . Other PBR parameters are also in the 0 – 1 range. The texture ID field is optional, and could also be null to indicate no texture.

Materials can be updated and are not immutable.

6.8. Lights. Lights are created, updated, and destroyed in the following messages.

LISTING 9. Light Messages

```

1 table LightCreateUpdate {
2     id          : LightID (required);
3     color       : Vec3;
4     intensity   : float;
5     // only point lights for now
6 }
7
8 table LightDelete {
9     id          : LightID (required);
10 }

```

The `color` key defines colors in the range of 0 – 1 for r, g, b . Intensity of the light is unbounded.

6.9. Signals & Methods. Signals and Methods are created, and destroyed in the following messages.

LISTING 10. Signals Messages

```

1 table MethodArg {
2     name : string;
3     doc  : string;
4 }
5
6 table MethodCreate {
7     id          : MethodID (required);
8     name        : string;
9     documentation : string;
10    returnDoc    : string;
11    argDoc       : [ MethodArg ];
12 }
13
14 table MethodDelete {
15     id          : MethodID (required);
16 }

```

Methods must be provided with a human friendly name. Two methods may not share the same name; there is no overloading. Documentation is recommended, but not required,

as is return value documentation. Argument information must be provided, at the very least given a name. You may not call a method with more arguments than as specified; use an argument that takes an array type to permit this option.

Signal must be provided with a human friendly name, and also may not share the same name. Arguments follow the same requirements as methods.

6.10. Signal Invoke & Method Reply. Signals may be invoked on the client and client methods replied to with the following messages.

LISTING 11. Communication

```

1  table SignalInvoke {
2      id      : SignalID (required);
3
4      // if the two below are not set, it is on the document
5      on_object : ObjectID;
6      on_table  : TableID;
7
8      signal_data : AnyList;
9  }
10
11 table MethodReply {
12     invoke_ident      : string (required);
13     method_data       : Any;
14     method_exception  : string;
15 }

```

Either only `on_object` or `on_table` or neither must be set, to indicate context. Signals may NOT be invoked on a context that does not have them attached.

Method replies must have a previously given method invocation identifier (see Section 7.3). If the method could not be executed, an exception field is filled instead of data. The method should let the caller know what went wrong, and who should fix it. Consider three cases:

Caller: The caller of the method (client) violated the contract of the method; incorrect number of arguments, bad argument content, etc.

Server/Application: A problem occurred with the application. The app itself encountered a problem; for example a client asking to load some data, but the server doesn't have permissions for that data. Or it could be a bug in the application itself.

Server/Library: The server library encountered an error.

6.11. Document. The document is implicit, and always exists. It can be modified with the following messages.

LISTING 12. Document Messages

```

1  table DocumentUpdate {
2      methods_list : [ MethodID ];
3      signals_list : [ SignalID ];
4  }

```

```

5 |
6 | table DocumentReset {
7 |     padding : bool; // these things cannot be empty, so...
8 | }

```

The document may be updated with `DocumentUpdate`, to modify the current methods and signals. It may also be completely reset. The reset message, by quirk of Flatbuffers, may not be empty; ignore any fields within. When a document is reset, all components and objects are to be deleted at that point.

7. CLIENT MESSAGES

In this section we discuss the messages sent by a client.

7.1. Root Message. Client messages are defined as the following root type.

LISTING 13. Client Message Root

```

1 | union ClientMessageType {
2 |     IntroductionMessage,
3 |     MethodInvokeMessage,
4 |     AssetRefreshMessage
5 | }
6 |
7 | table ClientMessage {
8 |     content : ClientMessageType (required);
9 | }
10 |
11 | table ClientMessages {
12 |     messages : [ ClientMessage ] (required);
13 | }
14 |
15 | root_type ClientMessages;

```

7.2. Introduction. The client introduces itself to the server with the following message.

LISTING 14. Introduction Message

```

1 | table IntroductionMessage {
2 |     client_name : string (required);
3 | }

```

The name of the client must not be empty, and should identify a client; host names can be used.

7.3. Method Invocation. The client asks to invoke a method with the following message.

LISTING 15. Method Invocation

```

1 | table MethodInvokeMessage {
2 |     methodId : MethodID (required);
3 |
4 |     // if the two below are not set, it is on the document

```

```

5 |     on_object : ObjectID;
6 |     on_table  : TableID;
7 |
8 |     invoke_ident : string (required);
9 |     method_args  : AnyList;
10| }

```

The message must have an invocation identifier; the asynchronous reply will carry that identifier. Identifiers must not be reused.

Either the `on_object` or the `on_table` or neither should be set, to indicate the context of the invocation: on an object, on a table, or on the document, respectively. The method can only be called on a context on which it is attached.

The arguments to the method must match the documented method signature.

7.4. Asset Refresh. The client may ask to receive missing buffer content with the following message.

LISTING 16. Buffer Refresh

```

1 | table AssetRefreshMessage {
2 |     for_buffers : [ BufferID ] (required);
3 | }

```

8. SEMANTICS

8.1. Tables. Tables are a way of exposing record data to clients so that they can either provide an alternative representation of that data or to allow command line clients access to the data. An example of an alternative representation would be a 2D chart that could be provided for a lightweight 2D client instead of a 3D plot.

Tables consist of columns (with unique names) and rows. Rows are identified by a `Key`, which is an integer. Keys are assumed to be monotonically increasing, that is, new insertions into the database are given a new key larger than any key seen before..

Another useful abstraction is the `Row` type; a row is either an `AnyList` or a `RealList`. A `Column` is the same.

A commonly used notion is the concept of a selection within a table of data. Listing 17 shows, in a JSON-like way, the definition of a Selection object as encoded in a NOODLES Any.

LISTING 17. Selection object definition. Note that the `to` field in the row ranges is exclusive. The `row_ranges` list *must* have an even number of elements.

```

1 | {
2 |     "rows" : [Key, ...],
3 |     "row_ranges" : [
4 |         Key from, Key to,
5 |         ...
6 |     ]

```

Method Name	Description
<code>void tbl_subscribe()</code>	Subscribe to changes in the table. The client will then receive signals, starting with the <code>table_init</code> signal.
<code>void tbl_insert_row(Row)</code>	Request to add data to the table.
<code>void tbl_insert_many([Column])</code>	Request to add many rows of data to the table, as a pack of column segments.
<code>void tbl_update(Key, Row)</code>	Request to update data to the table.
<code>void tbl_update_n([Key], [Column])</code>	Request to update many rows of data to the table, as a pack of column segments.
<code>void tbl_remove([Key])</code>	Ask to remove a list of keys.
<code>void tbl_clear()</code>	Ask to remove all rows of the table.
<code>void tbl_update_selection(/*snip*/) </code>	Ask to update a selection in the table.

TABLE 3. Table Methods summary

Signal Name	Description
<code>void tbl_init(/*snip*/) </code>	Initialize the table. Sent when the <code>subscribe</code> method is called.
<code>void tbl_updated([Key], [Column])</code>	Rows were updated in the table.
<code>void tbl_rows_removed([Key])</code>	Rows in the table were removed.
<code>void tbl_selection_updated(/*snip*/) </code>	A selection has changed.

TABLE 4. Table Signals summary

7 | }

8.1.1. *Methods & Signals.* To query table information, signals and methods are used. These names are restricted and cannot be used by the user application. Note, indexes are all zero-based. Tables 3 and 4 list the data related methods and signals a table can support. The server should not send any data or signals to the client for a given table *unless* a client has expressed interest by calling the `subscribe` method. This is to avoid stressing clients that have no table interface and to reduce unnecessary network traffic. Further it is up to the server to honor these methods; should the server not support modification, for example, requests will be dropped.

Subscribe. This allows the client to receive signals from the table. Without this, no signal should be sent by the server regarding the table. When this call is made, the server should reply with a `table_init` signal. The full signature of the signal is as follows:

```
void tbl_init(
    [ string ], // 1
```



```
[ Key ], //2
[ Column ], // 2
[ [string, SelectionObject] ] // 3
)
```

Argument 1 is a list of columns, including the key column as the first. This establishes a column order that is used to interpret and pack data values later in other calls and signals. The second argument provides the current data that is in the table, including the key column. The third is a pack of the current selections that are available in the table; this is provided as a list of pairs, where the first part of the pair is the string identifier of the selection and the second is the selection object that defines the selection.

Should the server re-issue the `table_init` signal, this would imply that the table has been reset/redefined.

Insertion. Data may be inserted into the table through both the row and many versions of the call. Note that the key cannot be specified, thus the row length should be for all the other columns, in column order. The many version simply takes a list of rows to be inserted. Insertion success is demonstrated through reception of the `rows_inserted` signal; this signal provides the data inserted along with the keys that were assigned to that row, i.e. the full row of data for all columns.

Update. Data can be updated through both the row and many versions. In this case, as opposed to the insertion functions, the full row, including the key column, is specified in column order, so that the correct row may be updated. Success is indicated through the corresponding update signal.

Removal. Data can be removed by specifying a list of keys to delete. Success will be indicated through the corresponding signal for all clients.

Selection. Data selections can be made through the `update_selection` call. The full signature of the call is as follows:

```
void tbl_update_selection( string, SelectionObject );
```

The first argument denotes the selection to update or add, and the selection object defines what that selection should be updated/initialized to. A selection object that is empty, i.e. specifying no rows or ranges is considered the empty selection and denotes that the selection should be deleted from clients.

This shall trigger the selection update signal. The full signature of the signal is as follows:

```
void tbl_selection_updated( string, SelectionObject );
```

This mirrors the update call, and denotes which selection has changed, and what to change it to.

8.1.2. Tables Metadata. Tables are also capable of synchronizing metadata for other purposes. This is exposed as a JSON object.

2D Plot Sync: To facilitate 2D plot synchronization, multiple optional mechanisms are present. The first provides a simple approach:

LISTING 18. Table Metadata for Plot Sync

```

1  Meta : {
2      "simple_plots" : [ SimplePlotInfo ] ,
3      <other keys>
4  }
5
6  SimplePlotInfo : {
7      "plot_name" : "name",
8      "column_name" : ColumnInfo,
9      ...
10 }
11
12 ColumnInfo : {
13     "prefers" : "x" | "y",
14     "color" : "#rrggbb",
15     "range" : [from, to]
16 }

```

In Listing 18, the metadata object will contain a key called `simple_plots`. This key is a listing of named plots; each plot describes how each column of the table should be treated in an arbitrary simple plot.

Complex 2D Plot Sync: More advanced plotting facilities are forthcoming, but planned to follow a system like: <http://docs.juliaplots.org/latest/attributes/>.

Direct Interface: Metadata may expose a `uri` field for direct database access. This is an object that contains information on how to connect.

LISTING 19. Table URI Field

```

1      Meta : {
2          "uri" : URI_Info ,
3          <other keys>
4      }
5
6      URI_Info : {
7          "address" : "ip",
8          "port" : number,
9          "type" : string
10     }

```

8.2. Objects. Objects may carry the logical operations.

For simplicity, in this section, we let `vec3 = RealList` and `vec4 = RealList`. When used as arguments, the three component and four component vectors require the exact number of components to be supplied in the list. Otherwise the server will consider that to be malformed, and can reject the call.

8.2.1. *Activator*. For clients, this could be when the user clicks on an object, or presses an interaction button when a wand is over an object.

```
void activate(string)
void activate(int)
list<string> get_activation_choices()
```

It is up to the server application to decide how to handle this ‘activation’. Activation is either in the string or integer form. Activation names can be obtained through the API (example: ‘Click’, ‘Clear Options’). An activation can be triggered by the string, or by an integer index. It makes sense to thus tie the order of names to priorities; a 0 is a primary click, 1 is an alternate click action, etc.

8.2.2. *Options*. Options are is conceptually the same thing as a combo-box widget.

```
list<string> get_option_choices()
string get_current_option()
void set_current_option(string)
```

A list of choices can be presented for an object, and an option can be set.

8.2.3. *Movable*. Movable objects allows the user to request to change the position of an object.

```
void set_position(vec3 p)
void set_rotation(vec4 q)
void set_scale(vec3 s)
```

Positions, rotations and scales are in the coordinate system of the parent object. The rotation is to be provided as a quaternion, with w being the last component.

8.2.4. *Selection*. Regions of an object can be ‘selected’. What this means is up to the application.

```
void select_region(vec3, vec3, bool)
void select_sphere(vec3, real, bool)
void select_half_plane(vec3, vec3, bool)
```

The selection API allows for a number of different selection tools. Others can be forged through the use of the movable API, and activators.

For **select_region**, the selection region is supplied as an axis aligned bounding box, and a boolean option for either selection (true) or deselection (false). For **select_sphere**, a position and a radius is supplied. For **select_half_plane**, a point and a normal is provided.

To support multiple selections, consider adding options and activators to your object.

8.2.5. *Query*. Objects can be probed to obtain a data value or annotation.

```
[string, vec3] probe_at(vec3)
```

The location (object local coordinates) to be probed is supplied in the argument. As a return value, a revised position is returned (in case the server desires to snap the probe to a different location) and a string containing the data to display.

Note that more complex actions may take place; a user can build their application to add more functionality (or use a different activator), which can instantiate objects for all users to see.

8.2.6. Annotation and Attention. The object may request user attention, through the following signals.

```
void signal_attention()
void signal_attention(vec3)
void signal_attention(vec3, string)
```

Multiple overloads are provided. If the signal has no data, the whole object would like attention. If there is a position, a specific object-local coordinate would like attention. If there is a string in addition to that, a message should be displayed at that point.

To attract attention, sounds, client-specific graphical adornment can all be used. For some clients, changing the camera view to include the point of attention can also be done.

8.2.7. Object Tags. Objects may be given tags. They are a list of strings. These allow the client to discover capabilities of the Object, or classify an object. Some tags imply the presence of certain methods or signals. Tags prefixed with `noo_` are reserved for use by the system.

Tag	Description
<code>noo_user_hidden</code>	On lists of objects or tree-views, this object should be hidden. Other objects should be visible ²

9. OPERATION & LIFECYCLE

9.1. Websocket Messages. The server side shall send the `ServerMessages` message, while clients are restricted to sending `ClientMessages` message.

9.2. Connection. Upon the connection of a websocket, the client first sends an introduction message. Any other message is ignored by the server until the introduction is provided.

The server will then send a list of creation messages to build the scene. This could pose a problem; large mesh or texture assets could take a significant amount of time to transfer and attempting to send those all at the start could cause issues ranging from the server being blocked, the client being overwhelmed, or de-synchronization, depending on implementations. In order to avoid this, the server may send creation info of buffers without the data. The client can use placeholder assets, and use the asset refresh mechanism to request asset updates with full information, which it can then use to update the graphical representation.

²This approach (hidden-specified) is chosen, because in a visible-specified, it is difficult to know when to hide the other objects.

From this point onward, the client can invoke methods, and the server can send signals and other messages.

APPENDIX A. COMMON FLATBUFFER SPECIFICATION

```
1 // Generate with
2 // flatc --scoped-enums --reflect-names --gen-mutable -c noodles.fbs
3
4 namespace noodles;
5
6 // Identifiers == These are tables, due to poor language support for
   structs.
7
8 table ObjectID {
9     id_slot : uint32;
10    id_gen  : uint32;
11 }
12
13 table TableID {
14     id_slot : uint32;
15     id_gen  : uint32;
16 }
17
18 table SignalID {
19     id_slot : uint32;
20     id_gen  : uint32;
21 }
22
23 table MethodID {
24     id_slot : uint32;
25     id_gen  : uint32;
26 }
27
28 table MaterialID {
29     id_slot : uint32;
30     id_gen  : uint32;
31 }
32
33 table GeometryID {
34     id_slot : uint32;
35     id_gen  : uint32;
36 }
37
38 table LightID {
39     id_slot : uint32;
40     id_gen  : uint32;
41 }
42
43 table TextureID {
44     id_slot : uint32;
```

```

45     id_gen  : uint32;
46 }
47
48 table BufferID {
49     id_slot : uint32;
50     id_gen  : uint32;
51 }
52
53 union AnyIDType {
54     ObjectID,
55     TableID,
56     SignalID,
57     MethodID,
58     MaterialID,
59     GeometryID,
60     LightID,
61     TextureID,
62     BufferID
63 }
64
65 table AnyID {
66     id : AnyIDType (required);
67 }
68
69 table MapEntry {
70     name  : string (key);
71     value : Any;
72 }
73
74 table Text { text : string; }
75 table Integer { integer : int64; }
76 table IntegerList { integers : [int64]; }
77 table Real { real : double; }
78 table RealList { reals : [double]; }
79 table Data { data : [byte]; }
80 table AnyList { list : [Any]; }
81 table AnyMap { entries : [MapEntry]; }
82
83 union AnyType {
84     Text,
85     Integer,
86     IntegerList,
87     Real,
88     RealList,
89     Data,
90     AnyList,
91     AnyMap,
92     AnyID
93 }
94

```

```

95 table Any {
96     any : AnyType (required);
97 }
98
99 // Misc Types
100     =====
101 struct Vec2 {
102     x : float;
103     y : float;
104 }
105
106 struct Vec3 {
107     x : float;
108     y : float;
109     z : float;
110 }
111
112 struct Vec4 {
113     x : float;
114     y : float;
115     z : float;
116     w : float;
117 }
118
119 struct Mat4 {
120     // for javascript compat, we have to expand this.
121     // components : [float : 16];
122     c1 : Vec4;
123     c2 : Vec4;
124     c3 : Vec4;
125     c4 : Vec4;
126 }
127
128 table BufferRef {
129     id      : BufferID;
130     start   : uint64;
131     size    : uint64;
132 }

```

APPENDIX B. SERVER MESSAGE FLATBUFFER SPECIFICATION

```

1 // Generate with
2 // flatc --scoped-enums --reflect-names --gen-mutable -c noodles_server.fbs
3
4 include "noodles.fbs";
5
6 namespace noodles;
7

```

```

8  // Server Messages
   =====
9
10 table MethodArg {
11     name : string;
12     doc  : string;
13 }
14
15 table MethodCreate {
16     id           : MethodID (required);
17     name         : string;
18     documentation : string;
19     returnDoc    : string;
20     argDoc       : [ MethodArg ];
21 }
22
23 table MethodDelete {
24     id           : MethodID (required);
25 }
26
27 //
   =====
28
29 table SignalCreate {
30     id           : SignalID (required);
31     name         : string;
32     documentation : string;
33     argDoc       : [ MethodArg ];
34 }
35
36 table SignalDelete {
37     id           : SignalID (required);
38 }
39
40 //
   =====
41
42 table TextDefinition { // Text plane, normal +z, up is +y, center: obj
   origin
43     text      : string (required); // String to render
44     font      : string (required); // Approximate font to use (e.g. Arial)
45     height    : float; // The height of the text. The width will be
46     opt_width : float; // Optional width of font.
47 }
48
49 // non-atomic update
50 table ObjectCreateUpdate {
51     id           : ObjectID (required);

```



```

52     name      : string;
53     parent    : ObjectID;
54     transform : Mat4;
55     material   : MaterialID;
56     mesh       : GeometryID;
57     lights     : [LightID];
58     tables     : [TableID];
59     instances  : [Mat4];
60     tags       : [string];
61     methods_list : [MethodID];
62     signals_list : [SignalID]; // Dont use "signals" to avoid Qt conflict.
63     text       : TextDefinition;
64 }
65
66 table ObjectDelete {
67     id      : ObjectID (required);
68 }
69
70 //
=====
71
72 table BufferCreate {
73     id      : BufferID (required);
74     bytes   : [byte]; // This could be empty
75     url     : string; // This could be empty too
76     url_size : uint64; // If you are passing a URL, this should be set
77
78     // If both are empty, data is coming, just hang on for another message
79     // of
80     // this kind with the bytes. you may have to request a refresh message
81 }
82
83 table BufferDelete {
84     id      : BufferID (required);
85 }
86 //
=====
87
88 table MaterialCreateUpdate {
89     id      : MaterialID (required);
90     color    : Vec4;
91     metallic : float;
92     roughness : float;
93     use_blending : bool;
94     texture_id : TextureID;
95 }
96

```

```

97 table MaterialDelete {
98     id          : MaterialID (required);
99 }
100
101 //
102     =====
103
104 table TextureCreateUpdate {
105     id          : TextureID (required);
106     reference   : BufferRef (required);
107 }
108
109 table TextureDelete {
110     id          : TextureID (required);
111 }
112 //
113     =====
114
115 table LightCreateUpdate {
116     id          : LightID (required);
117     color       : Vec3;
118     intensity   : float;
119     // only point lights for now
120 }
121
122 table LightDelete {
123     id          : LightID (required);
124 }
125 //
126     =====
127
128 table ComponentRef {
129     id          : BufferID (required);
130     start       : uint64;
131     size        : uint64;
132     stride      : uint64;
133 }
134
135 table GeometryCreate {
136     id          : GeometryID (required);
137
138     min_extent  : Vec3;
139     max_extent  : Vec3;
140
141     positions   : ComponentRef; // Vec3

```

```

141     normals      : ComponentRef; // Vec3
142     texCoords    : ComponentRef; // Vec2
143     colors       : ComponentRef; // U8Vec4
144     lines        : ComponentRef; // U16Vec2
145     triangles    : ComponentRef; // U16Vec3
146 }
147
148 table GeometryDelete {
149     id : GeometryID (required);
150 }
151
152 //
153     =====
154 // non-atomic update
155 table TableCreateUpdate {
156     id          : TableID (required);
157     name        : string;
158     meta        : string;
159     methods_list : [ MethodID ];
160     signals_list : [ SignalID ];
161 }
162
163 table TableDelete {
164     id          : TableID (required);
165 }
166
167 //
168     =====
169 table DocumentUpdate {
170     methods_list : [ MethodID ];
171     signals_list : [ SignalID ];
172 }
173
174 table DocumentReset {
175     padding : bool; // these things cannot be empty, so...
176 }
177
178 //
179     =====
180 table SignalInvoke {
181     id : SignalID (required);
182
183     // if the two below are not set, it is on the document
184     on_object : ObjectID;

```

```

185     on_table   : TableID;
186
187     signal_data : AnyList;
188 }
189
190 table MethodReply {
191     invoke_ident    : string (required);
192     method_data     : Any;
193     method_exception : string;
194 }
195
196 //
197
198 union ServerMessageType {
199     MethodCreate,
200     MethodDelete,
201     SignalCreate,
202     SignalDelete,
203     ObjectCreateUpdate,
204     ObjectDelete,
205     BufferCreate,
206     BufferDelete,
207     MaterialCreateUpdate,
208     MaterialDelete,
209     TextureCreateUpdate,
210     TextureDelete,
211     LightCreateUpdate,
212     LightDelete,
213     GeometryCreate,
214     GeometryDelete,
215     TableCreateUpdate,
216     TableDelete,
217     DocumentUpdate,
218     DocumentReset,
219     SignalInvoke,
220     MethodReply
221 }
222
223 table ServerMessage {
224     message : ServerMessageType;
225 }
226
227 table ServerMessages {
228     messages : [ ServerMessage ];
229 }
230
231 root_type ServerMessages;

```

APPENDIX C. CLIENT MESSAGE FLATBUFFER SPECIFICATION

```

1  // Generate with
2  // flatc --scoped-enums --reflect-names --gen-mutable -c noodles_client.fbs
3
4  include "noodles.fbs";
5
6  namespace noodles;
7
8  // Client Messages
9  =====
10 table IntroductionMessage {
11     client_name : string (required);
12 }
13
14 table MethodInvocationMessage {
15     methodId : MethodID (required);
16
17     // if the two below are not set, it is on the document
18     on_object : ObjectID;
19     on_table : TableID;
20
21     invoke_ident : string (required);
22     method_args : AnyList;
23 }
24
25 table AssetRefreshMessage {
26     for_buffers : [ BufferID ] (required);
27 }
28
29 union ClientMessageType {
30     IntroductionMessage,
31     MethodInvocationMessage,
32     AssetRefreshMessage
33 }
34
35 table ClientMessage {
36     content : ClientMessageType (required);
37 }
38
39 table ClientMessages {
40     messages : [ ClientMessage ] (required);
41 }
42
43 root_type ClientMessages;

```